

PaperReading

MIC-SVM:Designing a highly efficient support vector machine for advanced modern multi-core and many-core architectures

Y. You et al., IPDPS 2014

Bo Peng

Digital Science Center
Indiana University

pengb@indiana.edu

April 13, 2017

Outline

1 Background of the paper

- the Authors

2 Work of the paper

- I. Introduction
- II. Background of SVM
- III. Methodology for MIC-SVM
- IV. Experimental results and analysis
- VI. Conclusion

3 Comments

- Pros and Cons

the Authors



- Yang You, PhD student at UC Berkeley CS, Master 2015 Tsinghua
- research interests include Parallel Computing, Distributed Systems, and Machine Learning
- CA-SVM:
Communication-Avoiding Support Vector Machines on Distributed Systems. **Best Paper IPDPS 2015**

the Authors



- Prof. James Demmel
Computer Science Division
Chair, and EECS Associate
Chair, Professor of Mathematics
and Computer Science
- Research Projects:
Communication-Avoiding
Algorithms
The Complexity of Accurate
Floating Point Computation
LAPACK, scaLAPACK, main
contributor

the Authors: BeBOP Group

BeBOP(Berkeley Benchmarking and OPtimization Group)

The BeBOP group is broadly interested in understanding software performance tuning issues, and the interaction or implications for hardware design. Among our general interests are

- the interaction between application software, compilers, and hardware managing trade-offs among the various measures of performance, such as speed, accuracy, power, storage, ...
- automating the performance tuning process, starting with the computational kernels which dominate application performance in scientific computing and information retrieval
- performance modeling and evaluation of future computer architectures

Outline

1 Background of the paper

- the Authors

2 Work of the paper

- I. Introduction
- II. Background of SVM
- III. Methodology for MIC-SVM
- IV. Experimental results and analysis
- VI. Conclusion

3 Comments

- Pros and Cons

Motivations

Mic-svm: Designing a highly efficient support vector machine for advanced modern multi-core and many-core architectures

Motivation

- svm is important...widely used...recently being used in HPC for performance and power... prediction
- challenge of scaling performance over tens or even hundreds of cores within a single chip
- there has been no efficient SVM tool designed for advanced x86 based multi- and many-core architectures

Problems to solve(...unwritten...)

- What's the problem? –svm train phase inefficient in x-core system
- What's the state-of-art solutions? – libsvm, gpusvm
- What's the deficiency of those solutions? –
- What's the **methodology** to investigate these performance issues? – bottleneck, profiling, bound analysis, ...
- What's the suitable **optimization techniques?**
- What's a good **mapping** from input data pattern to specific architecture and optimization decisions?

Contributions

Contributions

- Designing and implementing MIC-SVM a highly efficient parallel...
- Proposing **novel analysis methods** and **optimization techniques** (i.e. adaptive support for input patterns and data parallelism) to fully utilize the multi-level parallelism provided by the studied architectures.
- Exploring and improving the deficiencies of the current tool...
- Providing insights on **how to select the most suitable architectures** for specific algorithms and input data patterns to achieve best performance.

Outline

1 Background of the paper

- the Authors

2 Work of the paper

- I. Introduction
- II. **Background of SVM**
- III. Methodology for MIC-SVM
- IV. Experimental results and analysis
- VI. Conclusion

3 Comments

- Pros and Cons

Classification Problems

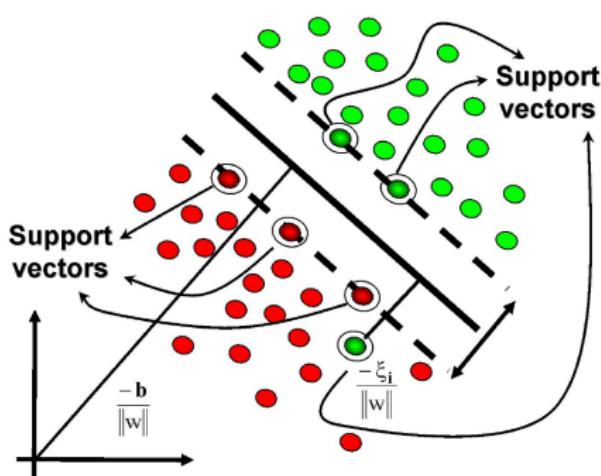


IRIS Dataset

| s.no | Sepal.Length | Sepal.Width | Petal.Length | Petal.Width | Species |
|------|--------------|-------------|--------------|-------------|-----------|
| 1 | 5.1 | 3.5 | 1.4 | 0.2 | setosa |
| 2 | 4.9 | 3.0 | 1.4 | 0.2 | setosa |
| 3 | 4.7 | 3.2 | 1.3 | 0.2 | setosa |
| 4 | 4.6 | 3.1 | 1.5 | 0.2 | setosa |
| 5 | 5.0 | 3.6 | 1.4 | 0.2 | setosa |
| 6 | 5.4 | 3.9 | 1.7 | 0.4 | setosa |
| s.no | Sepal.Length | Sepal.Width | Petal.Length | Petal.Width | Species |
| 145 | 6.7 | 3.3 | 5.7 | 2.5 | virginica |
| 146 | 6.7 | 3.0 | 5.2 | 2.3 | virginica |
| 147 | 6.3 | 2.5 | 5.0 | 1.9 | virginica |
| 148 | 6.5 | 3.0 | 5.2 | 2.0 | virginica |
| 149 | 6.2 | 3.4 | 5.4 | 2.3 | virginica |
| 150 | 5.9 | 3.0 | 5.1 | 1.8 | virginica |

- features: {sepal.length, sepal.width, petal.length, petal.width}.
- labels: {setosa, virginica}
- task: predict label by observed data with the fetures

SVM(Support Vector Machine)



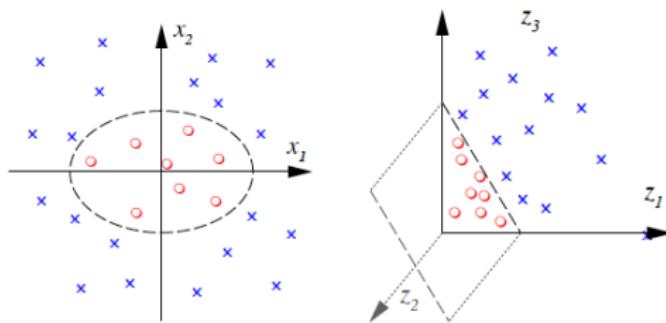
- separation plane:
 $y = \vec{w} \cdot \vec{x} - b$
- separation with maximum margin
- finally, we will get:
 $\vec{w} = \sum_{i=1}^n \alpha_i y_i \vec{x}_i$
 \vec{w} can be written as a linear combination of the **support vectors**.

- a representation of the examples as points in space, mapped so that the examples of the separate categories are divided by a clear gap that is **as wide as possible**.
- $\arg \min_{\vec{w}, b} \left\{ \frac{1}{2} \|\vec{w}\|^2 \right\}$ subject to $y_i(\vec{w} \cdot \vec{x}_i - b) \geq 1, \forall i$
- $\arg \min_{\vec{w}, b, \xi} \left\{ \frac{1}{2} \|\vec{w}\|^2 + C \sum_{i=1}^n \xi_i \right\}$ subject to $y_i(\vec{w} \cdot \vec{x}_i - b) \geq 1 - \xi_i, \forall i$

SVM: Kernel Trick

$$\Phi : \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

$$(x_1, x_2) \mapsto (z_1, z_2, z_3) := (x_1^2, \sqrt{2}x_1x_2, x_2^2)$$



$$\begin{aligned} k(\mathbf{x}, \mathbf{x}') &= \langle \mathbf{x}, \mathbf{x}' \rangle^2 \\ &= (x_1 x'_1 + x_2 x'_2)^2 \\ &= x_1^2 x'^2_1 + \sqrt{2} x_1 x_2 \sqrt{2} x'_1 x'_2 \\ &\quad + x_2^2 x'^2_2 \\ &= \langle \varphi(\mathbf{x}), \varphi(\mathbf{x}') \rangle \end{aligned}$$

- $\hat{y} = \vec{w} \cdot \vec{x} - b = \sum_{i=1}^n \alpha_i y_i k(\mathbf{x}_i, \mathbf{x}') - b$, for $k(\mathbf{x}, \mathbf{x}') = \langle \varphi(\mathbf{x}), \varphi(\mathbf{x}') \rangle_V$
- use of kernel functions, which enable them to operate in a high-dimensional, implicit feature space without ever computing the coordinates of the data in that space, but rather by simply computing the **inner products** between the images of all pairs of data in the feature space.

Numerical Optimization Problem

.....transformed into the dual form

- QP(Quadratic programming) problem
- training phase

Maximize: $F(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j K(X_i, X_j)$
subject to: $\sum_{i=1}^n \alpha_i y_i = 0$ and $0 \leq \alpha_i \leq C, \forall i$

- prediction phase

predict: $\hat{y}_i = u_i = \sum_{j=1}^n \alpha_j y_j K(\vec{x}_j, \vec{x}_i) - b$

SMO(Sequential Minimal Optimization) Algorithm

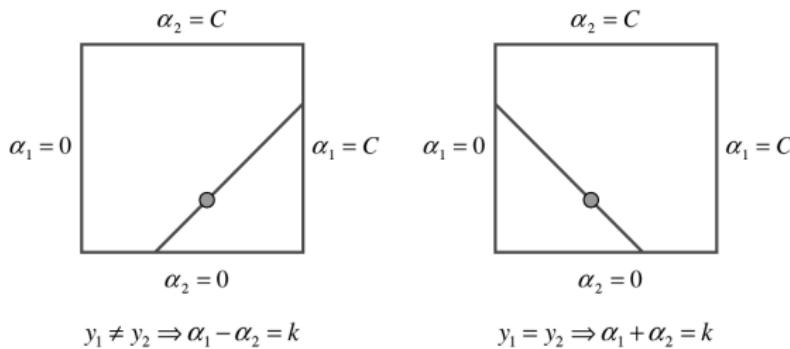
- At every step, SMO chooses **two** Lagrange multipliers to jointly optimize, finds the optimal values for these multipliers, and updates the SVM to reflect the new optimal values

subject to:

$$\sum_{i=1}^n \alpha_i y_i = 0$$

and

$$0 \leq \alpha_i \leq C, \forall i$$

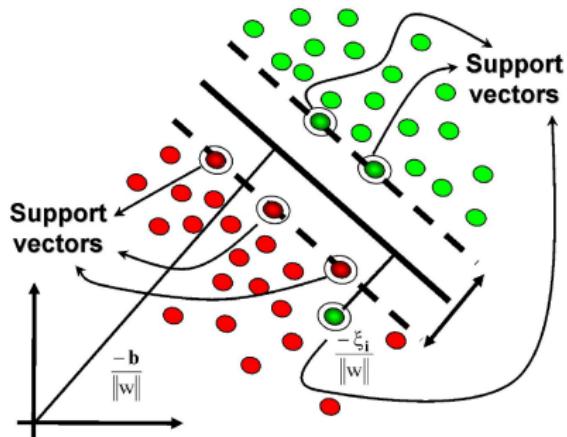


- The inequality constraints cause the Lagrange multipliers to lie in the box. The linear equality constraint causes them to lie on a diagonal line. Therefore, one step of SMO must find an optimum of the objective function on **a diagonal line segment**.

SMO: Update Equation

- The second derivative of the objective function along the diagonal line can be expressed as: $\eta = K(\vec{x}_1, \vec{x}_1) + K(\vec{x}_2, \vec{x}_2) - 2K(\vec{x}_1, \vec{x}_2)$
- $\alpha_2^{new} = \alpha_2 + \frac{y_2(E_1 - E_2)}{\eta}$ where $E_i = u_i - y_i$ is the error on ith training example
- $\alpha_1^{new} = \alpha_1 + y_1 y_2 (\alpha_2^{new} - \alpha_2)$
- $E_i^{new} = E_i + (\alpha_1^{new} - \alpha_1)y_1 K(\vec{X}_1, \vec{X}_i) + (\alpha_2^{new} - \alpha_2)y_2 K(\vec{X}_2, \vec{X}_i)$
- more details refer to [Platt 1998]

SMO:selection of the two lagrange multipliers



The Karush-Kuhn-Tucker (KKT) conditions are **necessary and sufficient conditions** for an optimal point of a positive definite QP problem.

set: $u_i = \sum_{j=1}^n \alpha_j y_j K(\vec{x}_j, \vec{x}_i) - b$,
then

$$\alpha_i = 0 \Leftrightarrow y_i u_i \geq 1,$$

$$0 < \alpha_i < C \Leftrightarrow y_i u_i = 1,$$

$$\alpha_i = C \Leftrightarrow y_i u_i \leq 1,$$

- chooses the first one by iterating over the entire training set, determining whether each example **violates the KKT conditions**.
- chooses the second one with **maximized step size** which is approximated by the absolute value of error $|E_1 - E_2|$.

SMO:selection equations

- (KKT) conditions

set: $u_i = \sum_{j=1}^n \alpha_j y_j K(\vec{x}_j, \vec{x}_i) - b,$

$\alpha_i = 0 \Leftrightarrow y_i u_i \geq 1,$

$0 < \alpha_i < C \Leftrightarrow y_i u_i = 1,$

$\alpha_i = C \Leftrightarrow y_i u_i \leq 1,$

- group the examples into two sets:

$I_1 = \{i : 0 < \alpha_i < C\} \cup \{i : y_i > 0, \alpha_i = 0\} \cup \{i : y_i < 0, \alpha_i = C\}$

where $E_i = u_i - y_i \geq 0$ otherwise violate KKT cond.

$I_2 = \{i : 0 < \alpha_i < C\} \cup \{i : y_i < 0, \alpha_i = 0\} \cup \{i : y_i > 0, \alpha_i = C\}$

where $E_i = u_i - y_i \leq 0$ otherwise violate KKT cond.

- selection equations:

$$i_1 = \arg \min_i \{E_i : i \in I_1\}$$

$$i_2 = \arg \max_i \{E_i : i \in I_2\}$$

SMO:Algorithm PseudoCode

Algorithm 1 The Original SMO Algorithm

0: Start

1: Input the training samples X_i and pattern labels y_i , $\forall i \in \{1, 2, \dots, n\}$.

2: Initializations, $\alpha_i = 0$, $f_i = -y_i$, $\forall i \in \{1, 2, \dots, n\}$.

3: Initializations, $b_{high} = -1$, $i_{high} = \min\{i : y_i = -1\}$, $b_{low} = 1$, $i_{low} = \max\{i : y_i = 1\}$.

4: Update α_{high} and α_{low} according to Equation (4) and (5).

5: If $Kernel(X_{high}, X_i)$ is not in memory, then compute $Kernel(X_{high}, X_i)$ and cache $Kernel(X_{high}, X_i)$ in memory using the LRU strategy.

6: If $Kernel(X_{low}, X_i)$ is not in memory, then compute $Kernel(X_{low}, X_i)$ and cache $Kernel(X_{low}, X_i)$ in memory through LRU strategy.

7: Update f_i according to Equation (6), $\forall i \in \{1, 2, \dots, n\}$

8: Compute i_{high} , i_{low} , b_{high} , and b_{low} according to Equation (9) and (10).

9: Update α_{high} and α_{low} according to Equation (4) and (5).

10: If iterations meet certain number, then do shrinking.

11: If $b_{low} > b_{high} + 2 \times tolerance$, then go to Step 5.

12: End

SMO:Summary

| Input Data | | Model | Cached |
|------------|-------|------------|--------|
| X1 | y1 | α_1 | E1 |
| X2 | y2 | α_2 | E2 |
| X3 | y3 | α_3 | E3 |
| X4 | y4 | α_4 | E4 |
| ... | ... | ... | ... |
| X_i | y_i | α_i | E_i |
| ... | ... | ... | ... |
| X_n | y_n | α_n | E_n |

- QP overall time complexity at least $O(n^2)$, space $O(n^2)$
- SMO type algorithm, empirically, it is known that the number of iterations may be higher than linear to the number of training data.
- $E_i^{new} = E_i + (\alpha_1^{new} - \alpha_1)y_1 K(\vec{X}_1, \vec{X}_i) + (\alpha_2^{new} - \alpha_2)y_2 K(\vec{X}_2, \vec{X}_i)$
- The update of all E_i at each step requires access to all the training samples, which costs more than 90% of the total time in the entire SMO algorithm.

Outline

1 Background of the paper

- the Authors

2 Work of the paper

- I. Introduction
- II. Background of SVM
- III. Methodology for MIC-SVM
- IV. Experimental results and analysis
- VI. Conclusion

3 Comments

- Pros and Cons

Methodology to parallelizing and optimizing

- In this section, we will briefly describe several important **design methods** and **optimization** details for our proposed MIC-SVM on parallelizing and accelerating SVM method on **Intel Ivy Bridge CPUs** and **Intel Xeon Phi coprocessor (MIC)**.

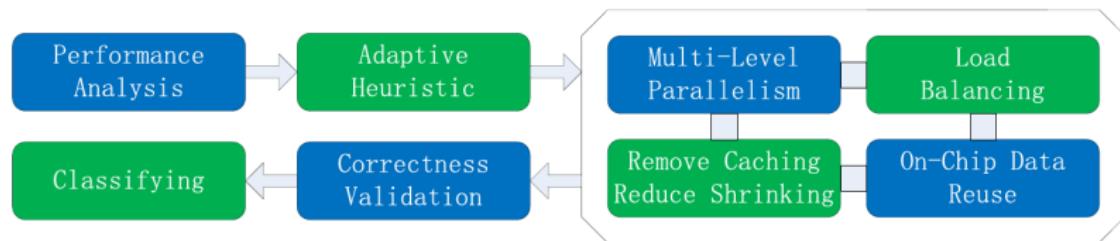


Fig. 2. General flow for parallelizing and optimizing MIC-SVM.

A: Analysis of the Algorithm and Architecture

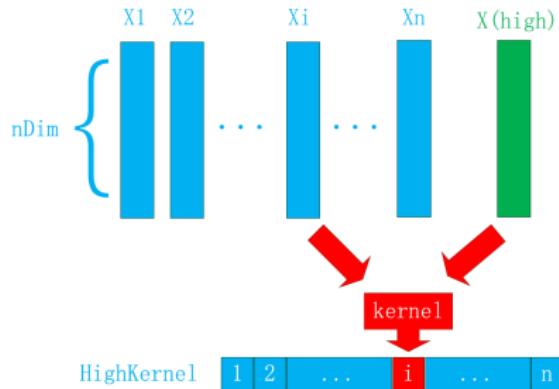


Fig. 3. This figure corresponds to the computation of HighKernel (used in Equation (11)). HighKernel is a vector, and $\text{HighKernel}(i) = \text{Kernel}(X_{\text{high}}, X_i)$. From this figure we can see that the evaluation of HighKernel need to access all the training samples. The computation of LowKernel can be accomplished in the same way.

- $E_i^{\text{new}} = E_i + (\alpha_1^{\text{new}} - \alpha_1)y_1 K(\vec{X}_1, \vec{X}_i) + (\alpha_2^{\text{new}} - \alpha_2)y_2 K(\vec{X}_2, \vec{X}_i)$
- The update of all E_i is the bottleneck.
- what's the bound?

A: Analysis of the Algorithm and Architecture

TABLE II
RCMB VALUES FOR EVALUATED ARCHITECTURES IN THIS PAPER

| Architecture | Ivy Bridge | MIC | Fermi | Kepler |
|-------------------------------------|------------|------|-------|--------|
| Frequency (GHz) | 2.20 | 1.09 | 1.50 | 0.73 |
| Peak Performance | | | | |
| Double (Gflops) | 422 | 1010 | 515 | 1320 |
| Peak Performance | | | | |
| Single (Gflops) | 844 | 2020 | 1030 | 3950 |
| Theoretical Memory bandwidth (GB/s) | 128 | 352 | 192 | 250 |
| RCMB | | | | |
| Single (Gflops/GB) | 6.59 | 5.74 | 5.36 | 15.8 |
| RCMB | | | | |
| Double (Gflops/GB) | 3.30 | 2.87 | 2.68 | 5.28 |

TABLE I
STANDARD KERNEL FUNCTIONS

| | |
|------------|--|
| Linear | $Kernel(X_i, X_j) = X_i \cdot X_j$ |
| Polynomial | $Kernel(X_i, X_j) = (aX_i \cdot X_j + r)^d$ |
| Gaussian | $Kernel(X_i, X_j) = -\gamma X_i - X_j ^2$ |
| Sigmoid | $Kernel(X_i, X_j) = \tanh(aX_i \cdot X_j + r)$ |

- Define the **Ratio of Computation to Memory Access** (RCMA) to describe SMOs algorithmic feature.

$$RCMA = \frac{\text{number_of_comp_flops}}{\text{number_of_memory_access_bytes}}$$

- Define Ratio of Peak Computation to Peak Memory Bandwidth (RCMB) to describe the theoretical architectural bound.

$$RCMB = \frac{\text{theoretical_peak_performance}}{\text{theoretical_bandwidth}}$$

C: Two-Level Parallelism

| arch | task parallelism | data parallelism |
|---------------|-------------------------------------|---|
| Ivy Bridge | utilizing multiple hardware threads | SIMD(AVX256) |
| MIC | utilizing multiple hardware threads | on-core VPU (Vector Processing Unit) and SIMD(AVX512) |
| Fermi, Kepler | independent warps | different CUDA cores |

C: Two-Level Parallelism

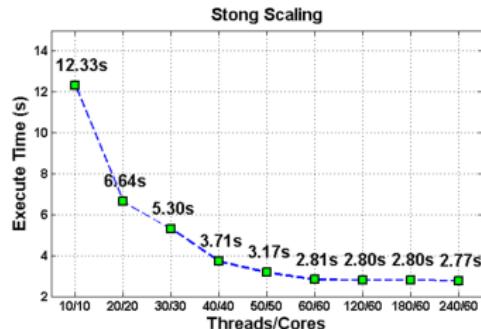


Fig. 6. Results for strong scaling test of MIC-SVM on Intel Xeon Phi using different number of threads and cores. threads/core=number of threads/number of cores.

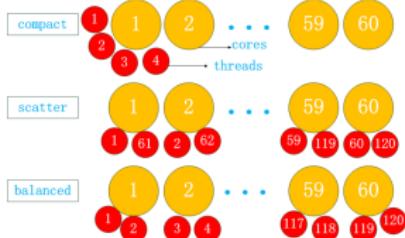


Fig. 7. Three affinity models for load balancing. 1) compact mode: the new

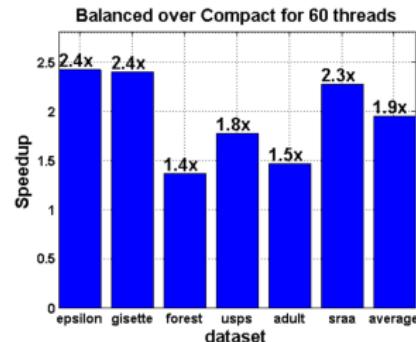


Fig. 8. Performance speedup of using "balanced mode" over "compact mode" when running MIC-SVM on Intel Xeon Phi with various types of datasets.

- The proper number of threads: number of physical resources in device
- Load balancing: affinity models
- use the Cilk array notation to achieve the data parallelism explicitly rather than applying the implicit compiler model.

D: Removing Caching and Reduce Shrinking Frequency

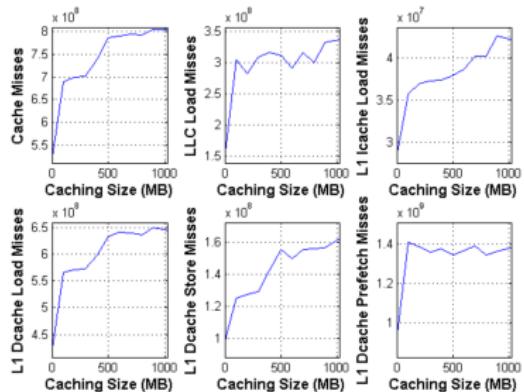


Fig. 9. Memory activities of running MIC-SVM on Intel Xeon Phi architecture while increasing the caching size .

- caching strategy can effectively reduce the number of kernels being evaluated. However, it also requires more memory access, may affect the performance negatively due to high memory access latency, memory contention from multithreading, and limited bandwidth.

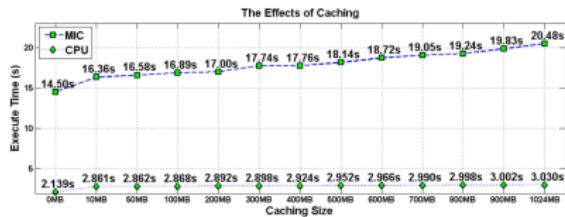


Fig. 10. Caching effects on overall execution time when running MIC-SVM on Intel Xeon Phi (MIC) and Ivy Bridge CPUs under increasing caching size.

E. Reducing the Gap Between RCMA and RCMB

- Together with load balancing, we employ the OpenMP SCHEDULE technique (static type, default chunk size) to **distribute the training samples evenly** to all hardware threads.
- exploring the proper **granularity of parallelism** by configuring the data sizes for two-level parallelism,
- minimizing the threads creation and destroy to reduce synchronization overhead
- maximizing the **TLB page size** to 2MB to obtain significantly more memory coverage when the datasets require more than 16MB memory.

Parallel SMO Algorithm

Algorithm 2 Parallel SMO Algorithm

0: Start

1: Input the training samples X_i and pattern labels y_i , $\forall i \in \{1, 2, \dots, n\}$. Applying adaptive heuristic support to decide how to process input datasets (sparse or dense) and whether to vectorize the kernel or not.

2: Map all the X_i , y_i , and α_i ($\forall i \in \{1, 2, \dots, n\}$) to all the hardware threads evenly through static SCHEME.

3: Initializations, for all hardware threads concurrently, $\vec{\alpha}_t = 0$, $\vec{F}_t = -\vec{Y}_t$, $\forall t \in \{1, 2, \dots, T\}$.

4: Initializations, $b_{high} = -1$, $i_{high} = \min\{i : y_i = -1\}$, $b_{low} = 1$, $i_{low} = \max\{i : y_i = 1\}$.

5: Check the thread Affinity to keep load balancing.

6: Update α_{high} and α_{low} according to Equation (4) and (5).

7: If kernel is vectorized, then do each kernel evaluation in each hardware thread through vectorization. Update \vec{F}_t according to Equation (6) through task parallelism, $\forall t \in \{1, 2, \dots, T\}$

8: If kernel is un-vectorized, then do each kernel evaluation in each hardware serially. Update \vec{F}_t according to Equation (6) through the combination of task parallelism and data parallelism, $\forall t \in \{1, 2, \dots, T\}$

9: Local reduce to obtain i_{high} , i_{low} , b_{high} , and b_{low} according to Equation (9) and (10) in each hardware thread.

10: Global reduce to obtain i_{high} , i_{low} , b_{high} , and b_{low} according to Equation (9) and (10).

11: Update α_{high} and α_{low} according to Equation (4) and (5).

12: If $b_{low} > b_{high} + 2 \times \text{tolerance}$, then go to Step 5.

13: End

Outline

1 Background of the paper

- the Authors

2 Work of the paper

- I. Introduction
- II. Background of SVM
- III. Methodology for MIC-SVM
- IV. Experimental results and analysis**
- VI. Conclusion

3 Comments

- Pros and Cons

A. Experimental Setup and Datasets

TABLE III
RELATED PARAMETERS OF ARCHITECTURES

| Architecture | Ivy Bridge | KNC | Fermi | Kepler |
|----------------------------------|------------|----------|----------|-----------|
| Cores | 24 | 61 | 448 | 2496 |
| L1 cache (KB) | 64/core | 64/core | 64/SM | 64/SM |
| L2 cache (KB) | 256/core | 512/core | 768/card | 1536/card |
| L3 cache (MB) | 30/socket | 0 | 0 | 0 |
| Coherent cache | L3 | L2 | L2 | L2 |
| Memory type | DDR3 | GDDR5 | GDDR5 | GDDR5 |
| Memory size (GB) | 64 | 8 | 6 | 6 |
| Measured Memory bandwidth (GB/s) | 68 | 159 | 97 | 188 |

TABLE IV
THE TEST DATASETS

| Dataset | <i>n</i> | <i>nDim</i> | Density | Cost | Gamma |
|--------------|----------|-------------|---------|-------|---------|
| epsilon [25] | 27,000 | 2,000 | 1.0000 | 1.0 | 0.08000 |
| gisette [26] | 6,000 | 5,000 | 0.9910 | 1.0 | 0.00020 |
| forest [27] | 12,000 | 54 | 0.2386 | 10000 | 0.00010 |
| usps [28] | 266,079 | 675 | 0.1497 | 1.0 | 0.03125 |
| adult [18] | 32,561 | 123 | 0.1128 | 0.1 | 0.06250 |
| sraa [21] | 72,309 | 20,958 | 0.0024 | 1.0 | 0.03125 |

C. Correctness Validation for MIC-SVM

TABLE V
THE CLASSIFICATION ACCURACY, THE VALUE OF B AND THE NUMBER OF SUPPORT VECTORS.

| Datasets | LIBSVM Accuracy-b-SVs | MIC-SVM on KNC MIC Accuracy-b-SVs | MIC-SVM on Ivy Bridge Accuracy-b-SVs | GPUSVM on Kepler Accuracy-b-SVs | GPUSVM on Fermi Accuracy-b-SVs |
|----------|--------------------------|--------------------------------------|---|------------------------------------|-----------------------------------|
| epsilon | 87.5%—0.0200—18116 | 87.5%—0.0200—18114 | 87.5%—0.0200—18113 | 87.5%—0.0200—18113 | 87.5%—0.0200—18113 |
| gisette | 97.7%—0.0034—1666 | 97.6%—0.0033—1665 | 97.6%—0.0033—1665 | 97.6%—0.0036—1665 | 97.6%—0.0036—1665 |
| forest | 82.2%—0.0120—11612 | 82.2%—0.0120—11609 | 82.2%—0.0120—11610 | 82.2%—0.0120—11610 | 82.2%—0.0120—11610 |
| usps | 99.2%—0.9464—39570 | 99.2%—0.9464—38596 | 99.2%—0.9464—38589 | 99.2%—0.9464—38581 | 99.2%—0.9464—38581 |
| adult | 84.4%—0.576—12281 | 84.4%—0.576—12273 | 84.4%—0.577—12273 | 84.4%—0.576—12278 | 84.4%—0.576—12278 |
| sraa | 97.0%—1.33—24430 | 97.0%—1.33—24392 | 97.0%—1.33—24395 | misconvergence | misconvergence |

- baseline: serial LIBSVM
- another solver on gpu: GPUSVM
- metrics: accuracy, support vectors number, b(discrepancy E)
- result: **nearly the same**

D. Performance Comparisons and Analysis

TABLE VI
ITERATIONS, TRAINING TIME AND SPEEDUPS.

| Datasets | LIBSVM Iterations-Time(s)-Speedup | MIC-SVM on MIC Iterations-Time(s)-Speedup | MIC-SVM on Ivy Bridge Iterations-Time(s)-Speedup | GPUSVM on Kepler Iterations-Time(s)-Speedup | GPUSVM on Fermi Iterations-Time(s)-Speedup |
|----------|--------------------------------------|--|---|--|---|
| epsilon | 11040–2959–1× | 11686–35.1–84× | 11616–64.2–46× | 11537–38.6–76× | 11537–46.1–64× |
| gisette | 1938–117.8–1× | 1887–2.75–43× | 1872–5.68–21× | 1894–5.46–22× | 1894–7.65–15× |
| forest | 23520–77.2–1× | 28912–17.6–4.4× | 28983–2.14–36× | 29018–2.04–38× | 28946–5.10–15× |
| usps | 34992–21945–1× | 47201–806–27× | 47173–884–25× | 47195–239–92× | 47195–429–51× |
| adult | 8028–84.1–1× | 8092–14.5–5.8× | 8150–2.14–39× | 8097–2.41–35× | 8097–5.67–15× |
| sraa | 14802–1128–1× | 13687–81.3–14× | 13676–24.9–45× | misconvergence | misconvergence |

- With powerful SIMD mechanism (512 bits), MIC is a good candidate for the **dense high-dimension datasets** with **modest number of training samples** because data parallelism can be achieved efficiently through vectorization;
- Equipped with sufficient caches and high clock frequency, Ivy Bridge CPUs are suitable for **sparse high-dimension datasets** since these datasets often require coarse-grained parallel processing;
- GPUs are proper for the datasets with **large number of training samples** and **low dimension** because they are more likely to benefit from millions of fine-grained lightweight threads.

D. Performance Comparisons and Analysis

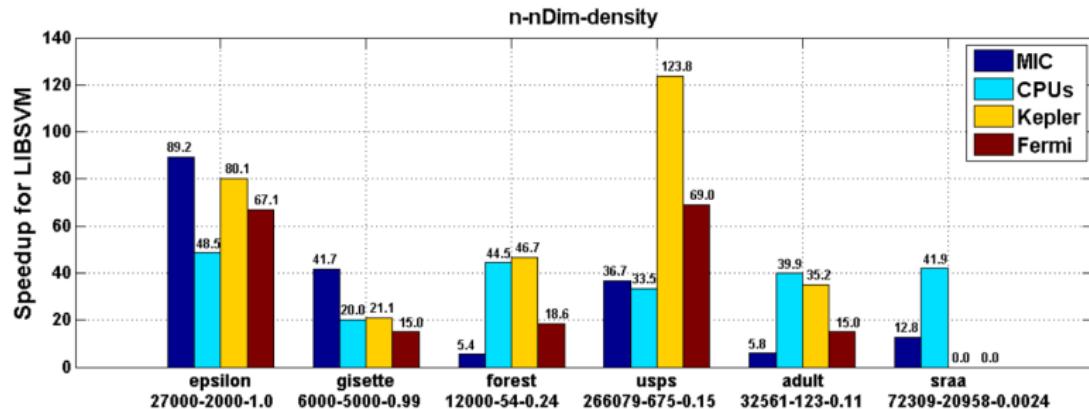


Fig. 11. This figure shows the speedups over LIBSVM based on the time of each iteration (rather than the whole time), n is the number of training samples, $nDim$ is the dimension of each training sample, and density represents the sparseness of a given dataset. All the results data shown in this figure are the average of many runs.

Outline

1 Background of the paper

- the Authors

2 Work of the paper

- I. Introduction
- II. Background of SVM
- III. Methodology for MIC-SVM
- IV. Experimental results and analysis
- VI. Conclusion

3 Comments

- Pros and Cons

Conclusion

Conclusion

- propose MIC-SVM, a highly efficient parallel support vector machine for x86 based multi-core and many-core architectures such as Intel Ivy Bridge CPUs and Intel KNC MIC.
- propose various novel analysis and optimization strategies that are general and can be easily applied to accelerate other machine learning methods.
- explore and improve the deficiencies of the current SVM tools.
- provide insights on how to map the most suitable architectures to specific data patterns in order to achieve the best performance.

Future work

- In future, we plan to extend the current MIC-SVM to distributed memory environment using multiple MICs.

Outline

1 Background of the paper

- the Authors

2 Work of the paper

- I. Introduction
- II. Background of SVM
- III. Methodology for MIC-SVM
- IV. Experimental results and analysis
- VI. Conclusion

3 Comments

- Pros and Cons

Comments on the paper

Pros

- gives a clear structure of the process of investigation.
- experiments are concise and informative.
- gives general analysis method and insights, which will be helpful to the community.

Cons

- well known optimizing techniques, and plain results, no much exciting things
- limited details on the effectiveness of optimizing techniques, e.g., the solution to fill the gap of RCMA and RCMB should be very cool, but...

References

-  Platt, J., 1998. Sequential minimal optimization: A fast algorithm for training support vector machines.
-  Y. You et al., Mic-svm: Designing a highly efficient support vector machine for advanced modern multi-core and many-core architectures, in Parallel and Distributed Processing Symposium, 2014 IEEE 28th International, 2014, pp. 809818.

The End