

# Parallel and Distributed GBT

9/18/2018

Bo Peng

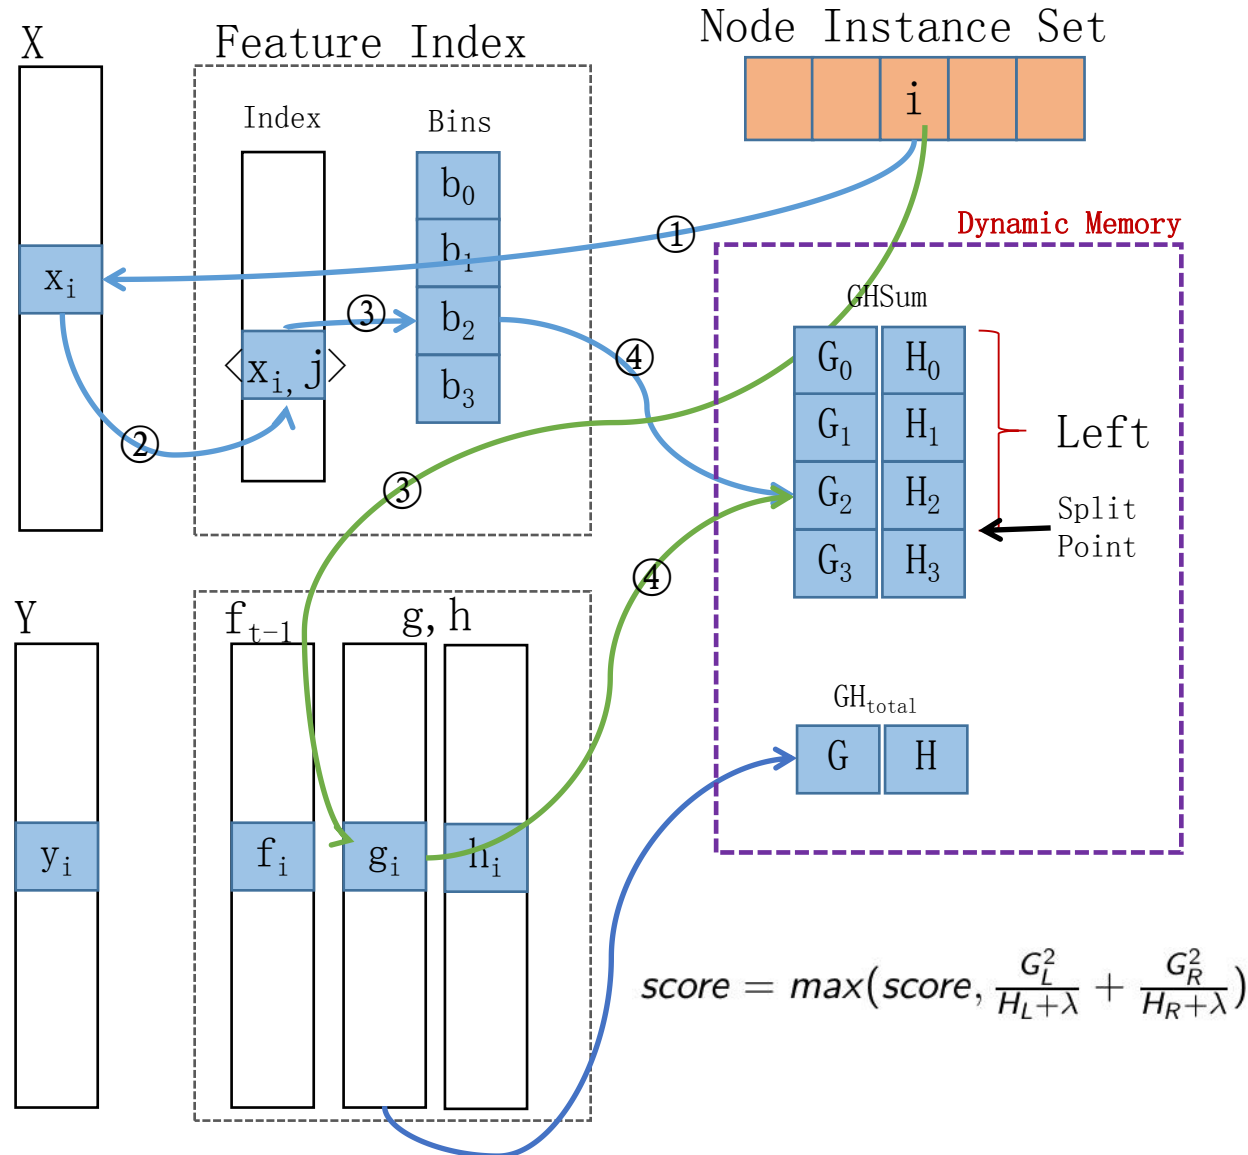
IPCC@Indiana University

# Outline

- Problem
- Related work
- Proposal



Single Feature

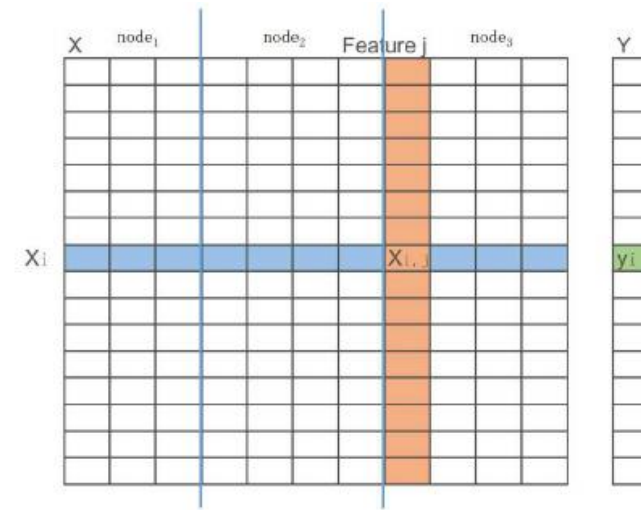


- `findBestSplit` is the bottleneck

- **GHSum** is the core data structure

- for each feature
- bins, fixed split points
- g,h summation on instances whose value fall into the bin

# Distributed Split



- Partition by rows (samples)
  - need global communication to build **Bins**, once if use global static Bins.
  - need global communication to build **GHSum** for each feature in `findBestSplit()` → `allreduce(GHSum)`
- Partition by columns (features)
  - **Bins** and **GHSum** are all local, no communication
  - need global communication to select the best feature in `findBestSplit()` → `allreduce(maxscore)`, also need to `broadcast(Split Instance Set)`

# Issues

- irregular memory access
  - instances in tree nodes are dynamic sets
  - non-continuous memory access to g,h
  - read/write after write dependency,  $GH[\text{bin}[\text{id}]] += g[\text{id}]$ , will stall when cache miss on access  $g[\text{id}]$
- sparsity
  - missing values
  - frequent zero entries
  - artifacts of feature engineering such as one-hot-encoding
- high dimensionality
  - $|\text{features}|$

# Outline

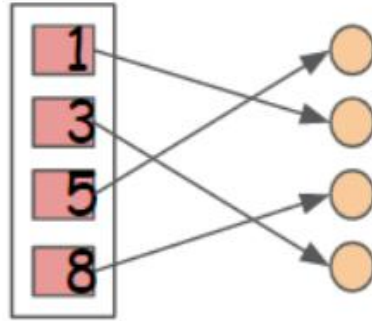
- Problem
- Related work
- Proposal

# XGBoost<sup>[1]</sup>

- Standard baseline
- Algorithm
  - approximate algorithm with both local and global proposal methods (**Bins**)
  - sparsity-aware, only collect statistics of non-missing entries
- System
  - **Column Block**
    - each block is a subset of rows
    - in each block, data stored in compressed column(CSC) format, with each column **sorted** by corresponding feature value
    - support feature level parallelism
    - linear scan to find best split
  - Cache-aware access
    - choose correct block size to get gradient statistics(g,h) fit into the CPU cache. ( $2^{16}$ )



Block Structure



Instructions

$G = G + g[\text{ptr}[i]]$   
 $H = H + h[\text{ptr}[i]]$

calculate score....

$G = G + g[\text{ptr}[i]]$   
 $H = H + h[\text{ptr}[i]]$



Short range  
dependency,  
with **non-contiguous**  
access to g

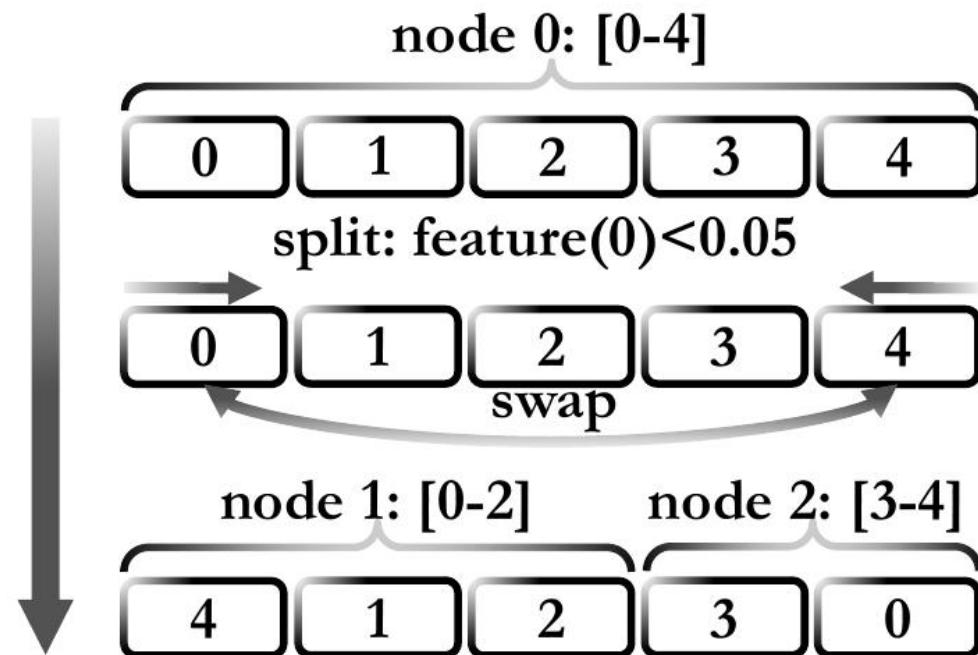
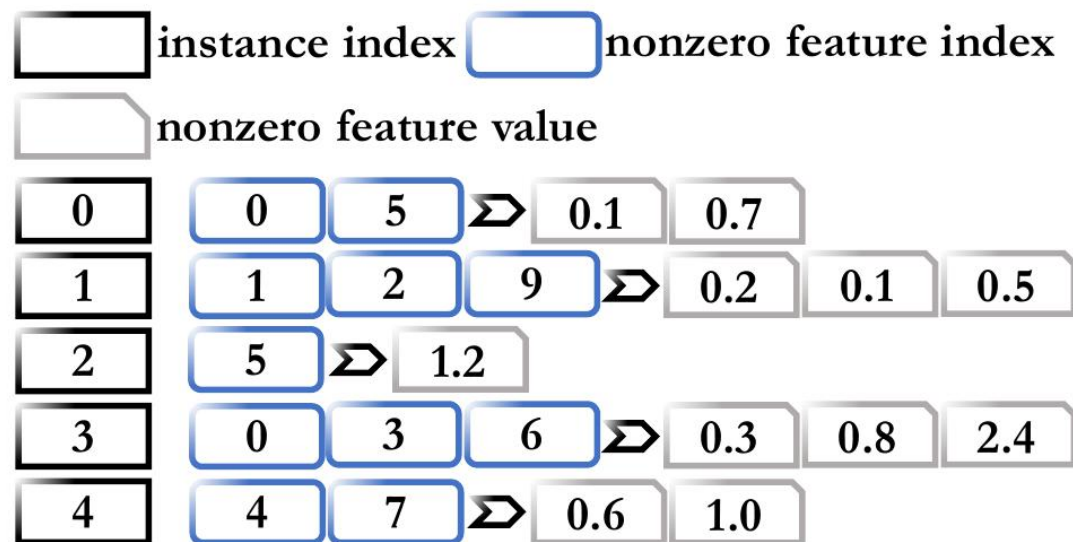
**Figure 8:** Short range data dependency pattern that can cause stall due to cache miss.

# LightGBM<sup>[2][3]</sup>

- high communication costs  $\sim O(|\text{features}| * \text{binSize})$ 
  - PV-Tree (Parallel Voting Decision Tree)
    - local voting: select top-k features based on local data
    - global voting: select top-2k features by votings from local candidates
    - collect full-grained histograms of the globally top-2k, and findBestSplit
- high dimensionality
  - Gradient-based One-side sampling
    - exclude a significant proportion of data instances with small gradients in estimate the score
    - better than SGB(Stochastic GB) with the same sampling ratio
- Sparsity
  - Exclusive feature bundling
    - bundle mutually exclusive features(never nonzero values simultaneously) into a single feature
    - letting exclusive features reside in different bins (adding offsets to original values)

# DimBoost<sup>[4][5]</sup>

- high dimensionality (industry application with 330K features)
- Optimize Communication
  - parameter server
    - claims to be one communication step and take less time (?)
  - two-phase-split finding
    - server-side split
  - round-robin task scheduler
    - schedule the splitting tasks among the workers
  - low-precision gradient histogram
    - each item  $q$  in a histogram, encode to a  $d$ -bit integer  $q' = \text{floor}(q/|c| * 2^d)$
    - $d=8$  often enough to obtain no loss on final accuracy
- Optimize Computation
  - parallel batch construction
    - divides a range into **batches** for the big nodes in the first few layers
  - sparse histogram construction
    - only non-zero entries



**Figure 9: Node-to-instance Index.**

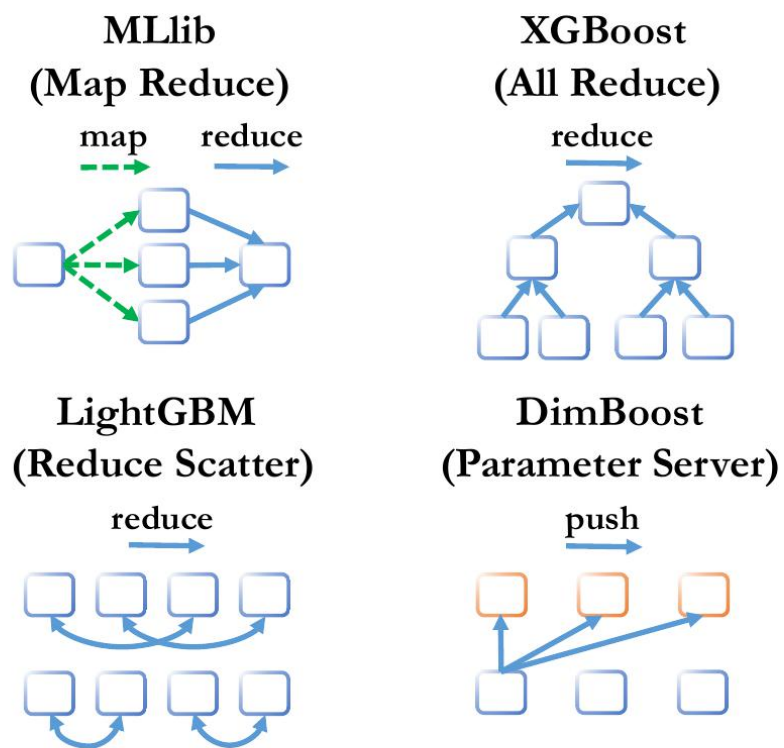


Figure 3: Existing Model Aggregation Methods.

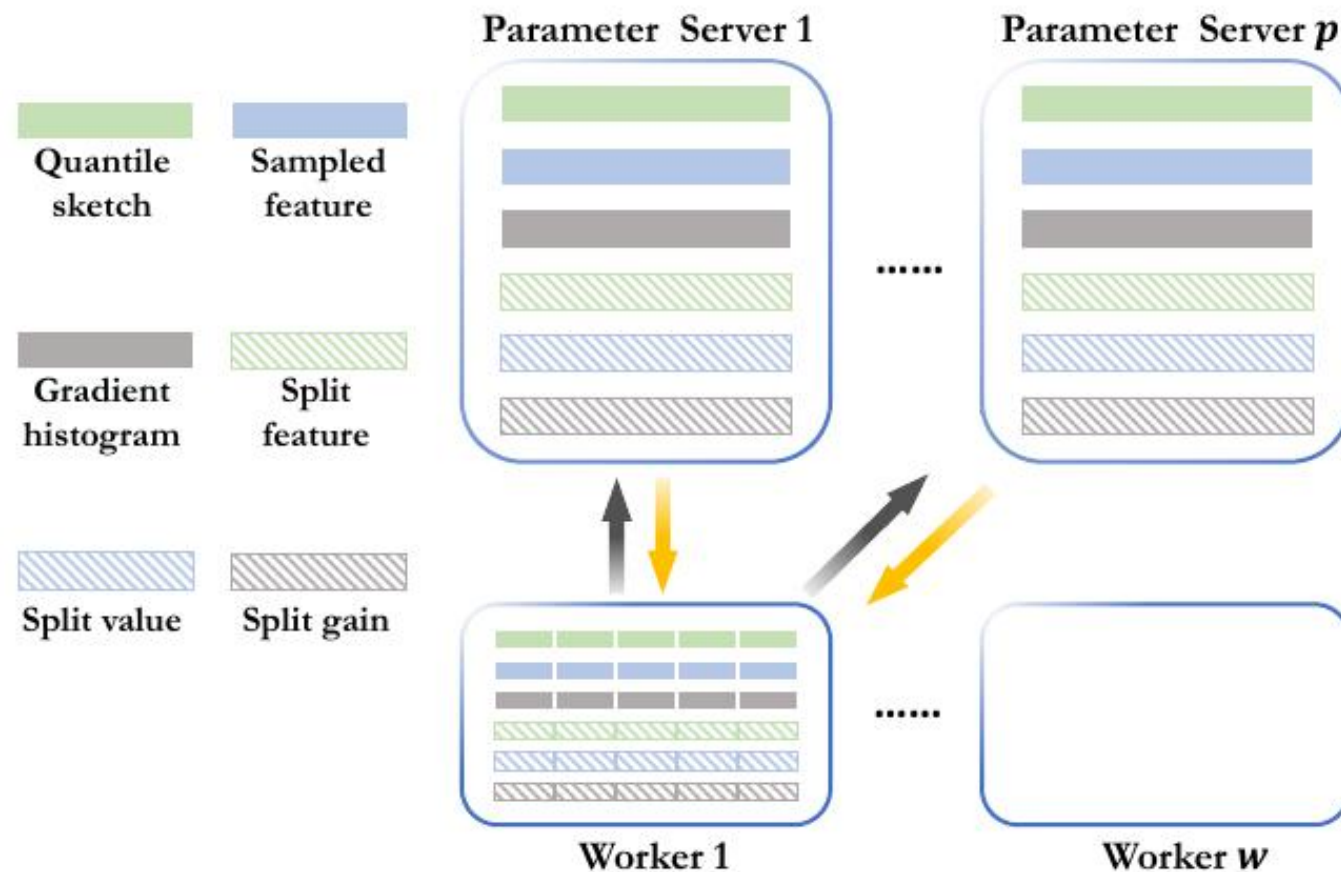


Figure 6: Parameter layout.

# GPU-GBDT<sup>[6]</sup>

- sparse representation
  - similar to Column Block
- **fine-grain** parallelism
  - parallelizing the score computation of each split point: segmented prefix sum  
--> GHSum
  - select best split point: segmented reduction
  - splitting a node: maintain values of each feature in the new node in sorted order
- Run-length encoding compression
  - sorted by feature values
  - repeated values can be compressed
  - directly splitting RLE elements
- thread/block workload dynamic allocation
  - under a memory constrain

TABLE I

DENSE AND SPARSE DATA REPRESENTATION

	Dense	Sparse
$\mathbf{x}_1$	$\langle 0.0, 0.0, 0.1, 0.0 \rangle$	$(a_3 : 0.1)$
$\mathbf{x}_2$	$\langle 1.2, 0.0, 0.1, 0.6 \rangle$	$(a_1 : 1.2); (a_3 : 0.1); (a_4 : 0.6)$
$\mathbf{x}_3$	$\langle 0.5, 1.0, 0.0, 0.0 \rangle$	$(a_1 : 0.5); (a_2 : 1.0)$
$\mathbf{x}_4$	$\langle 1.2, 0.0, 2.0, 0.0 \rangle$	$(a_1 : 1.2); (a_3 : 2.0)$

$$a_1 = (\mathbf{x}_2 : 1.2); (\mathbf{x}_4 : 1.2); (\mathbf{x}_3 : 0.5)$$

$$a_2 = (\mathbf{x}_3 : 1.0)$$

$$a_3 = (\mathbf{x}_4 : 2.0); (\mathbf{x}_2 : 0.1); (\mathbf{x}_1 : 0.1)$$

$$a_4 = (\mathbf{x}_2 : 0.6)$$

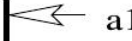


	 a1		 a2			
Instance Id	x1	x2	x4	x1	x3	x2
Gradient	0.7	-0.3	0.7	0.7	0.6	-0.3
<p>Segmented prefix sum</p> 						
Prefix sum	0.7	0.4	0.7	1.4	2	1.7

Fig. 1. Example results of segmented prefix sum

# Outline

- Problem
- Related work
- Proposal



# Assumption

- Assumption:
  - Model: GHSum, (Bins, g, h)
  - Model size  $\ll$  Training Data size
- Training Data  $N$  samples,  $M$  Features, build tree with  $L$  levels
- Model GHSum with bin size  $B$ 
  - at each node, size:  $B * M$
  - at each level, size:  $2^{L-1} * B * M$
- Most cases
  - $M < 1K$ ,  $L < 8$ ,  $B \sim [20, 100]$
  - high dimension can be an issue
  - deep tree can be an issue (GBT prefers shallow trees)

# Assumption

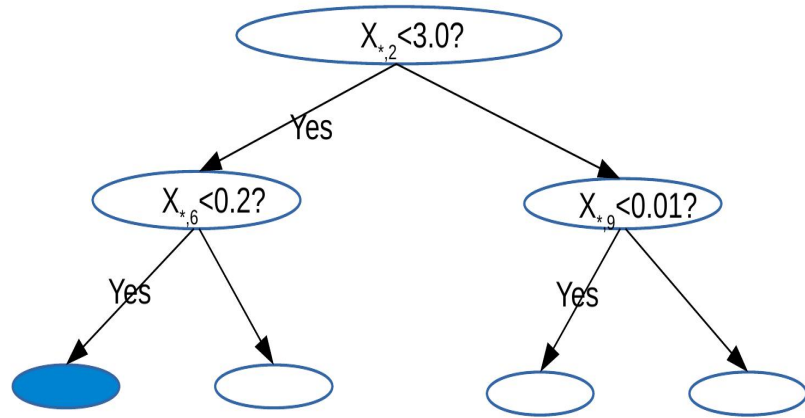
- example:
  - Higgs:  $N = 10\text{m}$ ,  $M = 28$ 
    - $L = 6$ ,  $B = 128$ ; model size at level  $L$ ,  $32 * 128 * 28 \ll 10\text{m}$
  - KDD10:  $N=19\text{M}$ ,  $M=29\text{M}$ 
    - feature engineering: expand a categorical feature into a set of binary features, combination\*

# Cache Friendly Data Organization



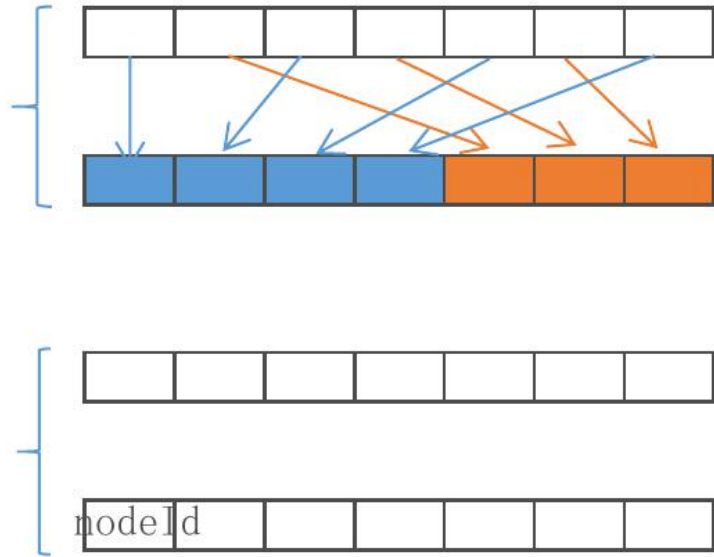
- Issue: non-continuous memory access in feature level parallelism
  - prefer column-wised data organization
  - dynamic instance set
  - N is large, (g,h) can not even fit into LLC cache
- Ideas
  - multi-columns to fit one **cache line**
  - **clolumn block** to get a small range of (g,h) access cache friendly

# Training data access



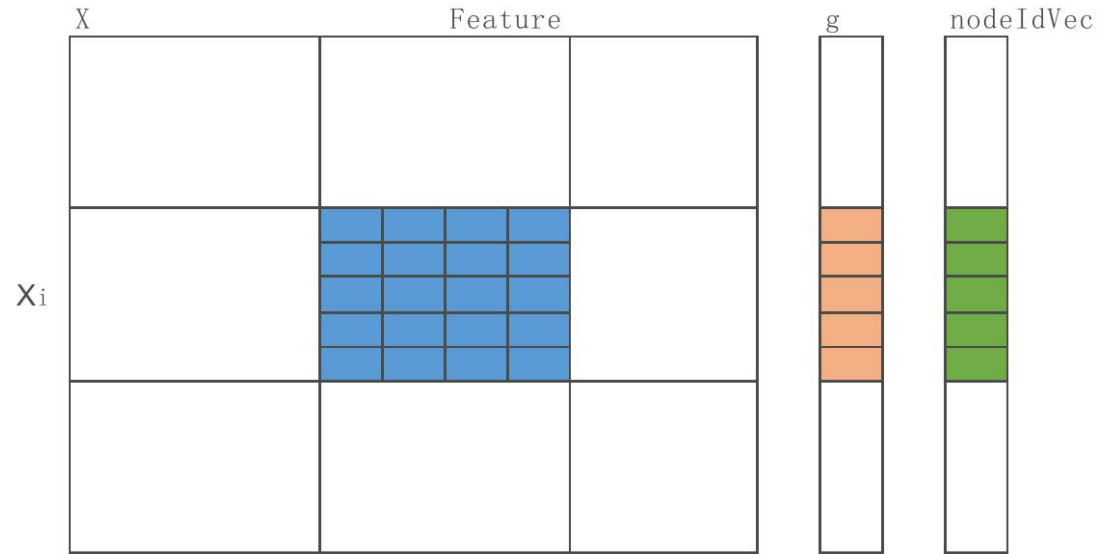
- Issue: redundant access to training data in node level parallelism
  - prefer row-wised data organization, however, it conflicts with feature level parallelism
- Ideas
  - Keep column-wised organization
  - Keep all GHSum of the same level in cache
  - **Single pass** on train data for each level. (half when reuse  $\text{GHSum}_{\text{parent}}$ )

# Sequential Scan on the Instance Set



- Issue: Dynamic instance set
  - move instance id around after node split
- Ideas
  - encode  $\langle \text{nodeid}, \text{level} \rangle$  into a unique nodeID, (easy for binary tree)
  - a nodeID vector to maintain the current leaf level instance sets directly
  - sequential access to this vector

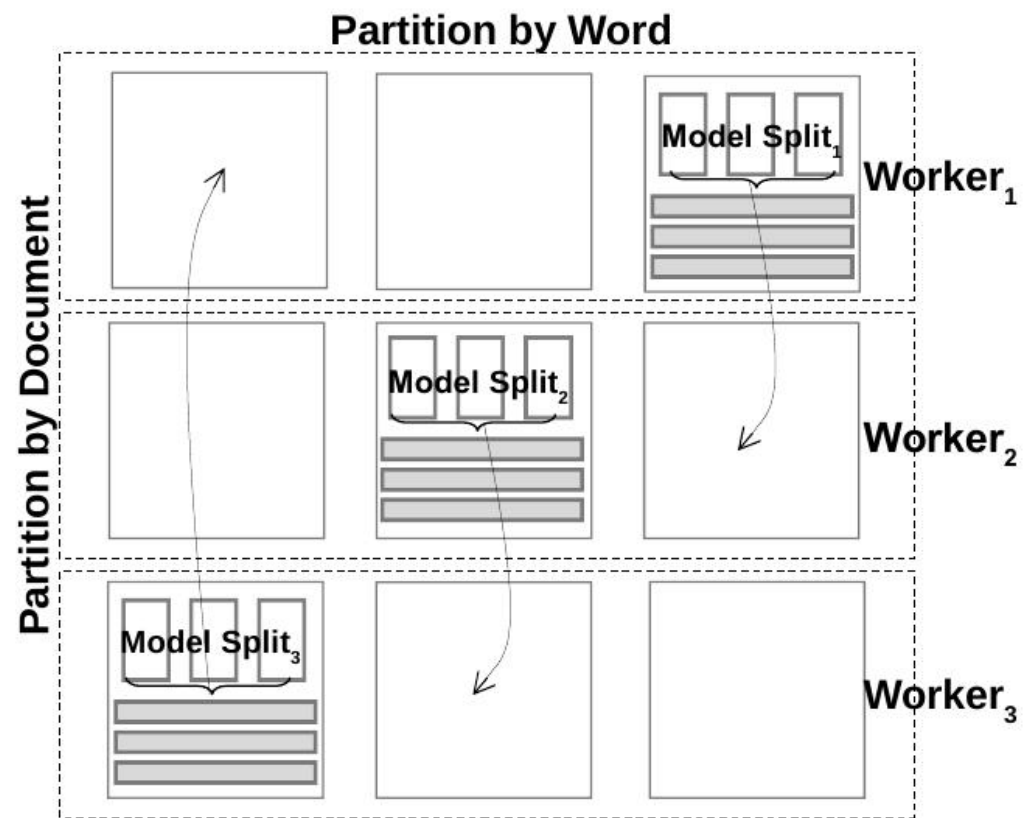
# Block-based Data Organization and Parallelism



```
for j in column_block:
    for each  $x_{i,j}$  in the column  $X_{,j}$ :
        nodeid = nodeIdVec[i]
        bindid = Bins[j][ $x_{i,j}$ ]
        GHSum[nodeid][j][bindid] += g[i]
```

- Column Block
  - the basic unit for parallelism
  - column-wised data organization
  - dense/sparse support
  - sorted by feature values
- Set the block number/size to
  - fit all active  $\langle g, h \rangle, \text{nodeIDVec}$  blocks in LLC
  - fit all active GHSum blocks in L2
  - loop order can be optimized

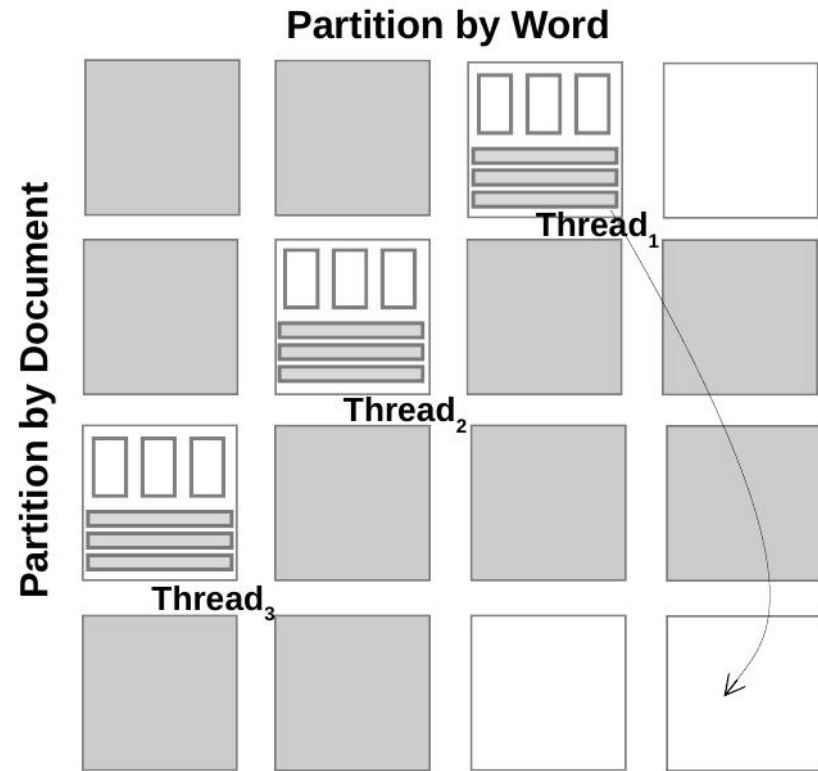
# Data and Model Parallelism by Model Rotation



(a) model rotation

- Proposal
  - use computation model : **model rotation**
- Data partitioned to blocks
  - Data Parallelism:
  - Model parallelism: GHSum split by feature columns
- Scheduler to avoid update conflicts

# Block Dynamic Scheduling in Shared Memory systems

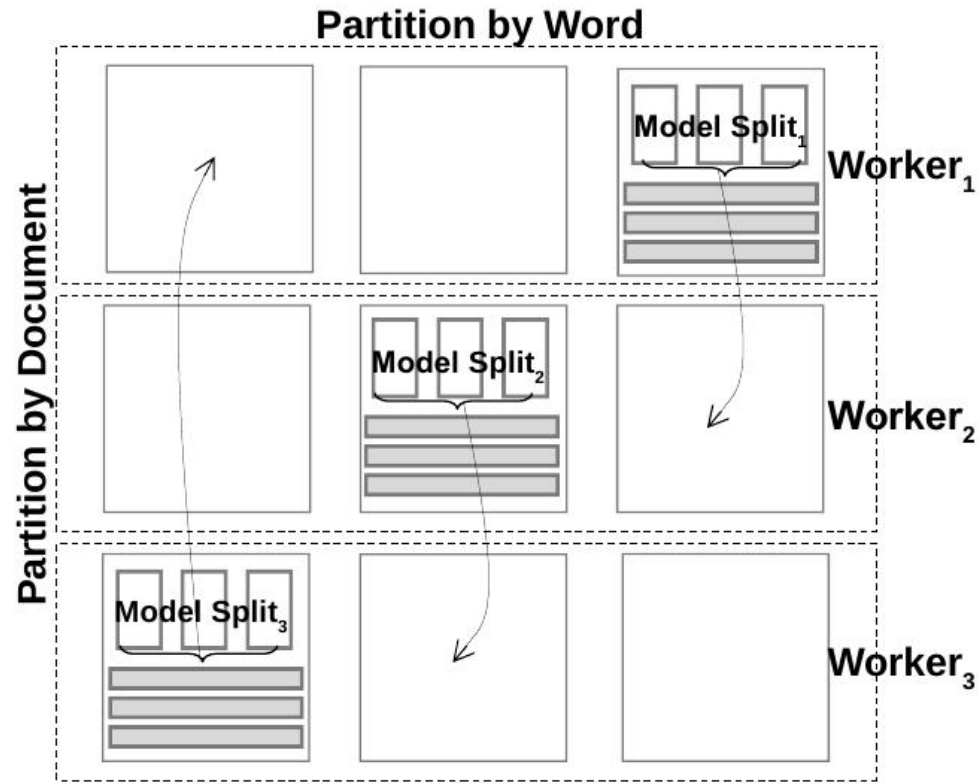


(b) dynamic scheduling

- a low cost solution to remove synchronization overhead
- number of partitions is larger than the number of threads,
- always ‘free’ rows and columns available when one thread finishes its current task.
- cache-aware block size
- sparse representation ready



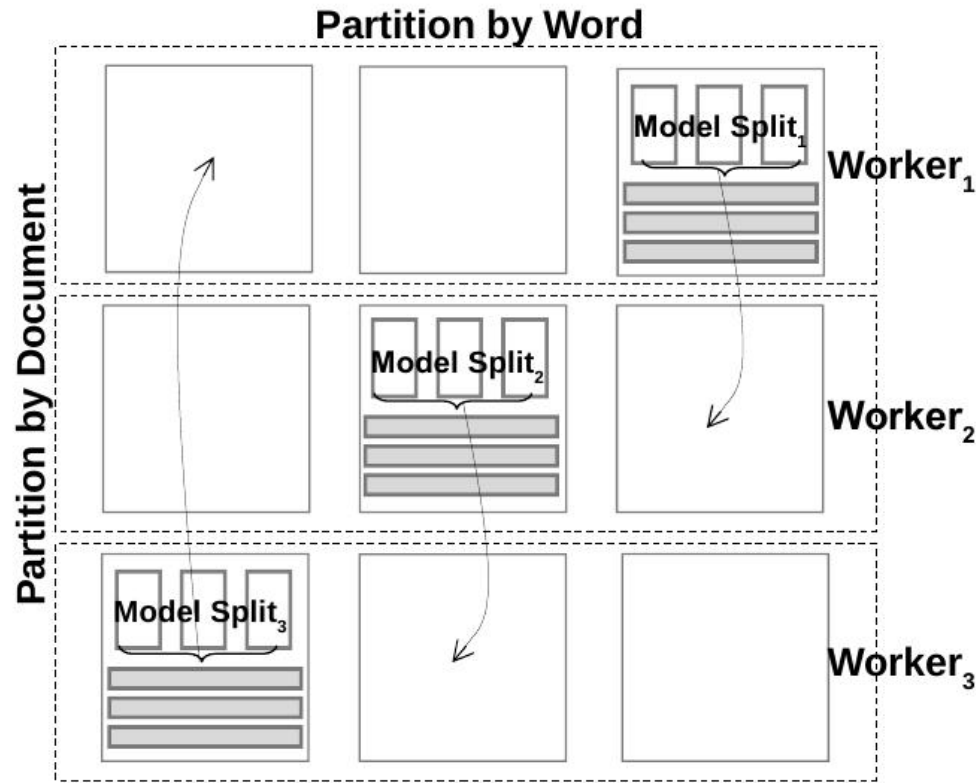
# Distributed Model Rotation



(a) model rotation

- Training data randomly partitioned among nodes
- Model partitioned among nodes
  - split by feature columns
- Scheduler to avoid update conflicts
  - rotate local model partition to neighbor at each sub-step
  - finish after K sub-steps(K nodes)
  - kind of a ring-allreduce

# Pipelining and Timer Control in Distributed Systems



(a) model rotation

- Each sampler only works for the same period of time and then the samplers do synchronization all together.
- They all use a **timer** to control the synchronization point rather than waiting until all the blocks to finish.
- This adjustment does not change the property of the uniform random selection of blocks.

# More to consider

- Feature Bundle
  - deal with category features
- Feature Sampling
  - design scheduler with Gradient-based priority
- Data Compression
  - model compression: low-precision compression
  - training data compression: RLE

# References

- [1]T. Chen and C. Guestrin, “Xgboost: A scalable tree boosting system,” in Proceedings of the 22nd acm sigkdd international conference on knowledge discovery and data mining, 2016, pp. 785–794.
- [2]Q. Meng et al., “A communication-efficient parallel algorithm for decision tree,” in Advances in Neural Information Processing Systems, 2016, pp. 1279–1287.
- [3]G. Ke et al., “LightGBM: A Highly Efficient Gradient Boosting Decision Tree,” in Advances in Neural Information Processing Systems, 2017, pp. 3149–3157.
- [4]J. Jiang, B. Cui, C. Zhang, and F. Fu, “DimBoost: Boosting Gradient Boosting Decision Tree to Higher Dimensions,” in Proceedings of the 2018 International Conference on Management of Data, 2018, pp. 1363–1376.
- [5]J. Jiang, J. Jiang, B. Cui, and C. Zhang, “TencentBoost: A Gradient Boosting Tree System with Parameter Server,” in 2017 IEEE 33rd International Conference on Data Engineering (ICDE), 2017, pp. 281–284.
- [6]Z. Wen, B. He, R. Kotagiri, S. Lu, and J. Shi, “Efficient Gradient Boosted Decision Tree Training on GPUs,” in 2018 IEEE International Parallel and Distributed Processing Symposium (IPDPS), 2018, pp. 234–243.